

Comprehensive Research: Databricks Apps, Databricks Lakehouse, and Databricks Asset Bundles

Research Date: October 20, 2025

Sources: Official Databricks Documentation, Databricks Community, GitHub Open Source Repositories

Table of Contents

1. [Executive Summary](#)
2. [Databricks Apps](#)
3. [Overview and Architecture](#)
4. [Key Concepts](#)
5. [Supported Frameworks](#)
6. [Development Workflow](#)
7. [Configuration and Deployment](#)
8. [Authentication and Security](#)
9. [Best Practices](#)
10. [Code Examples](#)
11. [Databricks Lakehouse](#)
12. [Architecture Overview](#)
13. [Core Components](#)
14. [Reference Architecture](#)
15. [Medallion Architecture](#)
16. [Databricks Lakebase \(OLTP\)](#)

17. [Implementation Guide](#)
 18. [Databricks Asset Bundles](#)
 19. [Overview and Concepts](#)
 20. [Bundle Structure](#)
 21. [Configuration Reference](#)
 22. [Development Lifecycle](#)
 23. [CI/CD Integration](#)
 24. [Deployment Modes](#)
 25. [Best Practices](#)
 26. [Code Examples](#)
 27. [Integration Patterns](#)
 28. [References](#)
-

Executive Summary

This comprehensive research document provides an in-depth analysis of three critical components of the Databricks Data Intelligence Platform: **Databricks Apps**, **Databricks Lakehouse**, and **Databricks Asset Bundles**. Each technology represents a fundamental capability for building modern data and AI applications on Databricks.

Databricks Apps enables developers to build and deploy production-ready data and AI applications directly within the Databricks platform, supporting popular frameworks like Streamlit, Dash, Flask, Gradio, and React. **Databricks Lakehouse** provides a unified architecture combining the best of data lakes and data warehouses, implementing the medallion architecture pattern for progressive data quality improvement. **Databricks Asset Bundles** offers infrastructure-as-code capabilities for managing Databricks resources, enabling robust CI/CD workflows and deployment automation.

This research synthesizes information from official Databricks documentation, community resources, and open-source GitHub repositories to provide actionable guidance for developers, data engineers, and architects.

Databricks Apps

Overview and Architecture

Databricks Apps is a production-ready platform for building, deploying, and hosting data and AI applications directly within the Databricks ecosystem. Released in October 2024, it provides a new modality for serving interactive applications that leverage the full power of the Databricks Data Intelligence Platform.

Core Value Proposition

Databricks Apps democratizes data intelligence by enabling even non-technical business analysts to access organizational data through intuitive, application-based interfaces. It eliminates the complexity of managing separate infrastructure for applications while maintaining enterprise-grade security, governance, and scalability.

Architecture Components

Databricks Apps operates as a **containerized service model** within the Databricks platform:

1. **Compute Layer:** Apps run on dedicated Databricks compute resources with configurable CPU and memory
2. **Runtime Environment:** Containerized execution environment supporting Python and Node.js runtimes
3. **Integration Layer:** Native integration with Databricks services (Unity Catalog, SQL Warehouses, Feature Store, Model Serving)
4. **Security Layer:** OAuth 2.0 authentication with dual identity model (app identity and user identity)
5. **Deployment Layer:** Automated deployment pipeline with version control and rollback capabilities

Key Concepts

App Structure

Every Databricks App consists of:

- **Source Code:** Application logic written in supported frameworks
- **Configuration File** (`app.yaml`): Defines runtime behavior, environment variables, and resource requirements
- **Dependencies:** Managed through `requirements.txt` (Python) or `package.json` (Node.js)
- **Static Assets:** Images, CSS, JavaScript files stored in the app directory
- **Compute Resources:** Configurable CPU/memory allocation

Dual Identity Model

Databricks Apps implements a sophisticated authentication model:

1. **App Identity:** Service principal or user account under which the app runs
2. **User Identity:** Individual user accessing the app, used for personalized data access

This enables apps to: - Access shared resources using the app identity - Enforce row-level security based on user identity - Maintain audit trails for compliance

Supported Frameworks

Python Frameworks

Framework	Type	Best For	Key Features
Streamlit	Data Apps	Rapid prototyping, dashboards	Simple API, reactive programming, built-in widgets
Dash	Analytics Apps	Complex dashboards, callbacks	Plotly integration, enterprise features
Gradio	ML Interfaces	Model demos, ML workflows	Auto-generated UI, easy sharing
Flask	Web Apps	Custom backends, APIs	Full control, lightweight, extensible
FastAPI	APIs	High-performance APIs	Async support, automatic documentation

Node.js Frameworks

Framework	Type	Best For	Key Features
React	Frontend	Interactive UIs	Component-based, virtual DOM, rich ecosystem
Angular	Frontend	Enterprise apps	Full framework, TypeScript, dependency injection
Svelte	Frontend	Lightweight apps	Compile-time optimization, minimal runtime
Express	Backend	APIs, middleware	Minimalist, flexible routing

Development Workflow

Step-by-Step Development Process

1. Environment Setup

```
# Install required dependencies
pip install gradio pandas databricks-sdk

# Create project directory
mkdir my-databricks-app
cd my-databricks-app
```

2. Create Application Code

Example Streamlit app (`app.py`):

```

import streamlit as st
import pandas as pd
from databricks import sql
import os

# Get environment variables set by Databricks
warehouse_id = os.getenv("DATABRICKS_WAREHOUSE_ID")
host = os.getenv("DATABRICKS_HOST")

st.title("Sales Analytics Dashboard")

# Connect to Databricks SQL Warehouse
@st.cache_resource
def get_connection():
    return sql.connect(
        server_hostname=host,
        http_path=f"/sql/1.0/warehouses/{warehouse_id}",
        credentials_provider=lambda: {} # Uses app identity
    )

# Query data
def load_data(query):
    conn = get_connection()
    cursor = conn.cursor()
    cursor.execute(query)
    return cursor.fetchall_arrow().to_pandas()

# UI Components
date_range = st.date_input("Select Date Range", [])
region = st.selectbox("Region", ["North", "South", "East", "West"])

if st.button("Load Data"):
    query = f"""
        SELECT date, region, SUM(sales) as total_sales
        FROM sales_data
        WHERE region = '{region}'
        GROUP BY date, region
        ORDER BY date
    """
    df = load_data(query)
    st.line_chart(df.set_index('date')['total_sales'])
    st.dataframe(df)

```

3. Configure App Runtime (app.yaml)

```

command: ['streamlit', 'run', 'app.py', '--server.port=8080']

env:
  - name: 'DATABRICKS_WAREHOUSE_ID'
    value: 'abc123def456'
  - name: 'STREAMLIT_GATHER_USAGE_STATS'
    value: 'false'
  - name: 'CATALOG_NAME'
    value: 'production'
  - name: 'SCHEMA_NAME'
    value: 'sales'

```

4. Define Dependencies (requirements.txt)

```
streamlit==1.28.0
pandas==2.1.0
databricks-sql-connector==3.0.0
plotly==5.17.0
```

5. Local Development and Testing

```
# Run locally
python app.py

# Or use Databricks CLI for local debugging
databricks apps run-local --prepare-environment --debug
```

6. Deploy to Databricks

```
# Deploy using Databricks CLI
databricks apps deploy my-app \
  --source-path . \
  --compute-size SMALL

# Or deploy via UI
# Navigate to Workspace → Apps → Create App
```

Configuration and Deployment

App.yaml Configuration Reference

The `app.yaml` file controls app execution behavior:

```

# Command to start the application
command:
  - gunicorn
  - app:app
  - -w
  - 4
  - --bind
  - 0.0.0.0:8080

# Environment variables
env:
  # Hardcoded values
  - name: 'APP_ENV'
    value: 'production'

  # Reference secrets from Databricks Secrets
  - name: 'API_KEY'
    valueFrom:
      secretKeyRef:
        scope: 'my-scope'
        key: 'api-key'

  # Reference Unity Catalog volumes
  - name: 'DATA_PATH'
    value: '/Volumes/catalog/schema/volume'

  # SQL Warehouse configuration
  - name: 'WAREHOUSE_ID'
    value: '${var.warehouse_id}'

```

Compute Size Configuration

Size	vCPUs	Memory	Best For
SMALL	2	8 GB	Development, low-traffic apps
MEDIUM	4	16 GB	Production apps, moderate traffic
LARGE	8	32 GB	High-traffic apps, complex processing
XLARGE	16	64 GB	Enterprise apps, heavy workloads

Deployment Logic

Databricks Apps uses intelligent deployment logic:

Default Behavior: - **Python apps:** Executes `python <first_py_file>` - **Node.js apps:** Executes `npm run start`

Custom Commands: Override defaults in `app.yaml` :


```

# Flask with Gunicorn
command: ['gunicorn', 'app:app', '-w', '4', '--bind', '0.0.0.0:8000']

# Streamlit with custom port
command: ['streamlit', 'run', 'app.py', '--server.port=8501']

# FastAPI with Uvicorn
command: ['uvicorn', 'main:app', '--host', '0.0.0.0', '--port', '8000']

```

Authentication and Security

OAuth 2.0 Integration

Databricks Apps implements enterprise-grade OAuth 2.0 authentication:

App-Level Authentication:

```

from databricks.sdk import WorkspaceClient

# Automatic authentication using app identity
w = WorkspaceClient()

# Access Unity Catalog
tables = w.tables.list(catalog_name="main", schema_name="default")

```

User-Level Authentication:

```

import os
from databricks.sdk.core import Config, oauth_service_principal

# Get current user context
user_email = os.getenv("DATABRICKS_USER_EMAIL")

# Enforce row-level security
query = f"""
    SELECT * FROM sensitive_data
    WHERE authorized_user = '{user_email}'
"""

```

Security Best Practices

- 1. Secret Management:** Store API keys and credentials in Databricks Secrets


```

` ``yaml env:
  - name: 'OPENAI_API_KEY' valueFrom: secretKeyRef: scope: 'ml-secrets' key: 'openai-key' `` `

```

2. **Unity Catalog Integration:** Leverage Unity Catalog for data governance


```
python
# Access governed data spark.sql("USE CATALOG production") df =
spark.table("customers") # Automatically enforces ACLs
```
3. **Network Security:** Apps run in isolated containers with controlled network access
4. **Audit Logging:** All app access and data queries are logged for compliance

Best Practices

Development Best Practices

1. **Modular Architecture**

```
my-app/ |— app.py # Main entry point |—
components/ # Reusable UI components | |— header.py | |— sidebar.py
| |— charts.py |— utils/ # Utility functions | |— data_loader.py |
|— auth.py |— assets/ # Static files | |— logo.png | |—
styles.css |— app.yaml # App configuration |— requirements.txt #
Dependencies
```

2. **Environment-Specific Configuration**

```
python import os
```

```
ENV = os.getenv("APP_ENV", "development")
```

```
if ENV == "production": DEBUG = False WAREHOUSE_ID =
os.getenv("PROD_WAREHOUSE_ID") else: DEBUG = True WAREHOUSE_ID =
os.getenv("DEV_WAREHOUSE_ID")
```

1. **Caching and Performance**

```
python import streamlit as st
```

```
@st.cache_data(ttl=3600) # Cache for 1 hour def load_large_dataset(): return
spark.table("large_table").toPandas()
```

```
@st.cache_resource def get_db_connection(): return sql.connect(...)
```

1. **Error Handling**

```
python try: data = load_data(query) except Exception as
e: st.error(f"Failed to load data: {str(e)}") st.stop()
```

Deployment Best Practices

1. **Use Databricks Asset Bundles** for managing apps as code

2. **Implement CI/CD pipelines** for automated testing and deployment
3. **Version Control:** Store app code in Git repositories
4. **Monitor Performance:** Use Databricks monitoring tools to track app metrics
5. **Implement Health Checks:** Add endpoints for monitoring app status

Code Examples

Example 1: Dash Analytics App

```
import dash
from dash import dcc, html, Input, Output
import plotly.express as px
from databricks import sql
import pandas as pd
import os

app = dash.Dash(__name__)

# Databricks connection
def get_data(query):
    connection = sql.connect(
        server_hostname=os.getenv("DATABRICKS_HOST"),
        http_path=f"/sql/1.0/warehouses/{os.getenv('WAREHOUSE_ID')}"
    )
    cursor = connection.cursor()
    cursor.execute(query)
    return cursor.fetchall_arrow().to_pandas()

# Layout
app.layout = html.Div([
    html.H1("Sales Performance Dashboard"),

    dcc.Dropdown(
        id='region-dropdown',
        options=[
            {'label': 'North America', 'value': 'NA'},
            {'label': 'Europe', 'value': 'EU'},
            {'label': 'Asia Pacific', 'value': 'APAC'}
        ],
        value='NA'
    ),

    dcc.Graph(id='sales-graph'),
    dcc.Graph(id='trend-graph')
])

# Callbacks
@app.callback(
    [Output('sales-graph', 'figure'),
     Output('trend-graph', 'figure')],
    [Input('region-dropdown', 'value')]
)
def update_graphs(region):
    query = f"""
        SELECT product, SUM(revenue) as total_revenue
        FROM sales
        WHERE region = '{region}'
        GROUP BY product
    """
    df = get_data(query)

    fig1 = px.bar(df, x='product', y='total_revenue',
                  title=f'Sales by Product - {region}')

    trend_query = f"""
        SELECT date, SUM(revenue) as daily_revenue
        FROM sales
        WHERE region = '{region}'
    """
```

```
        GROUP BY date
        ORDER BY date
    """
    df_trend = get_data(trend_query)
    fig2 = px.line(df_trend, x='date', y='daily_revenue',
                  title='Revenue Trend')

    return fig1, fig2

if __name__ == '__main__':
    app.run_server(host='0.0.0.0', port=8080)
```

Example 2: Flask API with ML Model

```
from flask import Flask, request, jsonify
from databricks.sdk import WorkspaceClient
import mlflow
import os

app = Flask(__name__)

# Load ML model from MLflow
w = WorkspaceClient()
model_name = os.getenv("MODEL_NAME")
model = mlflow.pyfunc.load_model(f"models://{model_name}/Production")

@app.route('/health', methods=['GET'])
def health():
    return jsonify({"status": "healthy"}), 200

@app.route('/predict', methods=['POST'])
def predict():
    try:
        data = request.get_json()
        features = data.get('features')

        # Make prediction
        prediction = model.predict([features])

        return jsonify({
            "prediction": prediction[0],
            "model_version": os.getenv("MODEL_VERSION")
        }), 200

    except Exception as e:
        return jsonify({"error": str(e)}), 500

@app.route('/batch-predict', methods=['POST'])
def batch_predict():
    try:
        data = request.get_json()
        table_name = data.get('table_name')

        # Load data from Unity Catalog
        df = spark.table(table_name)

        # Batch prediction
        predictions = model.predict(df)

        # Save results
        result_table = f"{table_name} predictions"
        predictions_df = df.withColumn("prediction", predictions)
        predictions_df.write.mode("overwrite").saveAsTable(result_table)

        return jsonify({
            "status": "success",
            "result_table": result_table,
            "rows_processed": df.count()
        }), 200

    except Exception as e:
        return jsonify({"error": str(e)}), 500

if __name__ == '__main__':
    app.run(host='0.0.0.0', port=8000)
```

Example 3: Gradio ML Interface

```
import gradio as gr
from databricks.sdk import WorkspaceClient
import mlflow
import pandas as pd

# Load model
model = mlflow.pyfunc.load_model("models:/customer-churn/Production")

def predict_churn(age, tenure, monthly_charges, total_charges):
    """Predict customer churn probability"""
    features = pd.DataFrame({
        'age': [age],
        'tenure': [tenure],
        'monthly_charges': [monthly_charges],
        'total_charges': [total_charges]
    })

    prediction = model.predict(features)[0]
    probability = model.predict_proba(features)[0][1]

    return {
        "Churn Prediction": "Yes" if prediction == 1 else "No",
        "Churn Probability": f"{probability:.2%}",
        "Retention Recommendation": get_recommendation(probability)
    }

def get_recommendation(probability):
    if probability > 0.7:
        return "High Risk - Immediate intervention required"
    elif probability > 0.4:
        return "Medium Risk - Monitor and engage"
    else:
        return "Low Risk - Standard retention program"

# Create Gradio interface
interface = gr.Interface(
    fn=predict_churn,
    inputs=[
        gr.Number(label="Customer Age"),
        gr.Number(label="Tenure (months)"),
        gr.Number(label="Monthly Charges ($)"),
        gr.Number(label="Total Charges ($)")
    ],
    outputs=gr.JSON(label="Prediction Results"),
    title="Customer Churn Prediction",
    description="Predict customer churn probability using ML model",
    examples=[
        [45, 24, 75.50, 1810.00],
        [32, 6, 120.00, 720.00],
        [58, 60, 55.25, 3315.00]
    ]
)

if __name__ == "__main__":
    interface.launch(server_name="0.0.0.0", server_port=7860)
```

Databricks Lakehouse

Architecture Overview

The **Databricks Lakehouse** represents a paradigm shift in data architecture, combining the scalability and flexibility of data lakes with the performance and ACID guarantees of data warehouses. It provides a unified platform for all data workloads—from ETL and BI to machine learning and real-time analytics.

Core Principles

1. **Unified Platform:** Single platform for all data workloads (batch, streaming, ML, BI)
2. **Open Standards:** Built on open formats (Delta Lake, Apache Iceberg, Parquet)
3. **ACID Transactions:** Full transactional guarantees for data reliability
4. **Schema Evolution:** Support for schema changes without breaking existing queries
5. **Time Travel:** Query historical versions of data for auditing and recovery
6. **Unified Governance:** Centralized governance through Unity Catalog

Core Components

1. Delta Lake

Delta Lake is the foundational storage layer providing ACID transactions on data lakes:


```

# Create Delta table
df.write.format("delta").mode("overwrite").save("/mnt/delta/events")

# Read Delta table
df = spark.read.format("delta").load("/mnt/delta/events")

# Time travel
df_yesterday = spark.read.format("delta") \
    .option("versionAsOf", 1) \
    .load("/mnt/delta/events")

# Update data with ACID guarantees
from delta.tables import DeltaTable

deltaTable = DeltaTable.forPath(spark, "/mnt/delta/events")
deltaTable.update(
    condition = "eventType = 'click'",
    set = { "processed": "true" }
)

# Merge (upsert) operation
deltaTable.alias("target").merge(
    source.alias("source"),
    "target.id = source.id"
).whenMatchedUpdate(set = {
    "value": "source.value",
    "updated_at": "current_timestamp()"
}).whenNotMatchedInsert(values = {
    "id": "source.id",
    "value": "source.value",
    "created_at": "current_timestamp()"
}).execute()

```

Key Features:

- **ACID Transactions:** Atomicity, Consistency, Isolation, Durability
- **Scalable Metadata:** Handles petabyte-scale tables efficiently
- **Time Travel:** Access historical data versions
- **Schema Enforcement:** Prevents bad data from corrupting tables
- **Audit History:** Complete history of all changes

2. Unity Catalog

Unity Catalog provides centralized governance for data and AI assets:

```

-- Create catalog
CREATE CATALOG production;

-- Create schema
CREATE SCHEMA production.sales;

-- Create managed table
CREATE TABLE production.sales.transactions (
  transaction_id STRING,
  customer_id STRING,
  amount DECIMAL(10,2),
  transaction_date DATE
) USING DELTA;

-- Grant permissions
GRANT SELECT ON TABLE production.sales.transactions TO `data_analysts`;
GRANT MODIFY ON TABLE production.sales.transactions TO `data_engineers`;

-- Row-level security
CREATE FUNCTION production.sales.filter_region(region STRING)
RETURNS BOOLEAN
RETURN current_user() IN (
  SELECT user_email FROM production.sales.regional_access
  WHERE allowed_region = region
);

ALTER TABLE production.sales.transactions
SET ROW FILTER production.sales.filter_region(region) ON (region);

```

Capabilities: - **Multi-cloud governance:** Works across AWS, Azure, GCP - **Fine-grained access control:** Table, column, and row-level security - **Data lineage:** Track data flow across transformations - **Audit logging:** Complete audit trail of data access - **Centralized metadata:** Single source of truth for all data assets

3. Apache Spark and Photon

Apache Spark provides distributed processing, while **Photon** is Databricks' vectorized query engine:

```

# Spark DataFrame operations
df = spark.read.table("production.sales.transactions")

# Complex transformations
result = df.filter(col("amount") > 1000) \
    .groupBy("customer_id") \
    .agg(
        sum("amount").alias("total_spent"),
        count("*").alias("transaction_count"),
        avg("amount").alias("avg_transaction")
    ) \
    .orderBy(desc("total_spent"))

# Write results
result.write.mode("overwrite").saveAsTable("production.sales.customer_summary")

# Photon automatically accelerates queries
spark.conf.set("spark.databricks.photon.enabled", "true")

```

Reference Architecture

The Databricks Lakehouse reference architecture consists of seven functional layers:

Layer 1: Source

Data originates from multiple sources:

- **Structured Sources:** Relational databases (PostgreSQL, MySQL, SQL Server, Oracle)
- **Semi-Structured:** JSON, XML, Avro, Parquet files
- **Unstructured:** Logs, images, videos, documents
- **Streaming:** Kafka, Kinesis, Event Hubs, IoT devices
- **SaaS Applications:** Salesforce, Workday, SAP via Lakeflow Connect

Layer 2: Ingest

Multiple ingestion patterns:

Batch Ingestion:

```
# Auto Loader for incremental file ingestion
df = spark.readStream.format("cloudFiles") \
    .option("cloudFiles.format", "json") \
    .option("cloudFiles.schemaLocation", "/mnt/schema/events") \
    .load("/mnt/landing/events/")

df.writeStream \
    .format("delta") \
    .option("checkpointLocation", "/mnt/checkpoints/events") \
    .table("bronze.events")
```

Streaming Ingestion:

```
# Kafka streaming
df = spark.readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "kafka:9092") \
    .option("subscribe", "transactions") \
    .load()

# Parse and write to Delta
parsed_df = df.selectExpr("CAST(value AS STRING) as json") \
    .select(from_json("json", schema).alias("data")) \
    .select("data.*")

parsed_df.writeStream \
    .format("delta") \
    .outputMode("append") \
    .option("checkpointLocation", "/mnt/checkpoints/kafka") \
    .table("bronze.transactions")
```

Lakeflow Connect (Built-in Connectors):

```
# Ingest from Salesforce
import dlt

@dlt.table
def salesforce_accounts():
    return spark.read.format("salesforce") \
        .option("sfObject", "Account") \
        .load()
```

Layer 3: Storage

Data stored in cloud object storage (S3, ADLS, GCS) using open formats:

- **Delta Lake tables:** Primary storage format
- **Apache Iceberg:** Alternative open table format
- **Parquet:** Columnar storage for analytics
- **Unity Catalog Volumes:** For unstructured data (PDFs, images, models)

Layer 4: Transform

Data transformation using multiple engines:

Delta Live Tables (Declarative ETL):

```
import dlt
from pyspark.sql.functions import *

@dlt.table(
    comment="Raw events from source systems",
    table_properties={"quality": "bronze"}
)
def bronze_events():
    return spark.readStream.table("source.events")

@dlt.table(
    comment="Cleaned and validated events",
    table_properties={"quality": "silver"}
)
@dlt.expect_or_drop("valid_timestamp", "timestamp IS NOT NULL")
@dlt.expect_or_drop("valid_user", "user_id IS NOT NULL")
def silver_events():
    return dlt.read_stream("bronze_events") \
        .withColumn("processed_at", current_timestamp()) \
        .dropDuplicates(["event_id"])

@dlt.table(
    comment="Aggregated user metrics",
    table_properties={"quality": "gold"}
)
def gold_user_metrics():
    return dlt.read("silver_events") \
        .groupBy("user_id", window("timestamp", "1 day")) \
        .agg(
            count("*").alias("event_count"),
            countDistinct("session_id").alias("session_count")
        )
```

Spark SQL:

```
-- Complex analytical query
WITH customer_segments AS (
    SELECT
        customer_id,
        SUM(amount) as total_spent,
        COUNT(*) as transaction_count,
        CASE
            WHEN SUM(amount) > 10000 THEN 'VIP'
            WHEN SUM(amount) > 5000 THEN 'Premium'
            ELSE 'Standard'
        END as segment
    FROM production.sales.transactions
    WHERE transaction_date >= '2024-01-01'
    GROUP BY customer_id
)
SELECT
    segment,
    COUNT(*) as customer_count,
    AVG(total_spent) as avg_spent,
    SUM(total_spent) as segment_revenue
FROM customer_segments
GROUP BY segment
ORDER BY segment_revenue DESC;
```

Layer 5: Query/Process

Multiple compute options for different workloads:

Compute Type	Best For	Characteristics
SQL Warehouses	BI, analytics, ad-hoc queries	Serverless, auto-scaling, optimized for SQL
All-Purpose Clusters	Interactive development, notebooks	Persistent, customizable, multi-language
Job Clusters	Scheduled ETL, batch processing	Ephemeral, cost-effective, isolated
Serverless Compute	On-demand workloads	Instant startup, pay-per-use

Layer 6: Serve

Data serving for different consumption patterns:

Data Warehousing:

```

-- Create materialized view for BI
CREATE MATERIALIZED VIEW production.analytics.sales_summary AS
SELECT
    DATE_TRUNC('month', transaction_date) as month,
    product_category,
    region,
    SUM(amount) as total_revenue,
    COUNT(DISTINCT customer_id) as unique_customers
FROM production.sales.transactions
GROUP BY month, product_category, region;

-- Optimize for query performance
OPTIMIZE production.analytics.sales_summary
ZORDER BY (month, product_category);

```

Model Serving:

```

# Deploy ML model for real-time serving
from databricks.sdk import WorkspaceClient
from databricks.sdk.service.serving import ServedEntityInput,
EndpointCoreConfigInput

w = WorkspaceClient()

w.serving_endpoints.create(
    name="customer-churn-model",
    config=EndpointCoreConfigInput(
        served_entities=[
            ServedEntityInput(
                entity_name="main.ml_models.customer_churn",
                entity_version="3",
                workload_size="Small",
                scale_to_zero_enabled=True
            )
        ]
    )
)

```

Lakebase (OLTP):

```

-- Create OLTP database instance
CREATE DATABASE INSTANCE my_oltp_db
WITH (
    instance_type = 'db.t3.medium',
    storage_size = 100
);

-- Create transactional table
CREATE TABLE my_oltp_db.orders (
    order_id SERIAL PRIMARY KEY,
    customer_id INTEGER NOT NULL,
    order_date TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    status VARCHAR(20),
    total_amount DECIMAL(10,2)
);

-- Sync to Delta Lake for analytics
CREATE SYNCED TABLE production.sales.orders_sync
AS SELECT * FROM my_oltp_db.orders;

```

Layer 7: Analysis

Final consumption layer:

- **BI Tools:** Tableau, Power BI, Looker connected via JDBC/ODBC
- **Databricks SQL Editor:** Native SQL interface with dashboards
- **Notebooks:** Interactive analysis with Python, R, Scala, SQL
- **Applications:** Databricks Apps for custom interfaces
- **APIs:** REST APIs for programmatic access

Medallion Architecture

The **Medallion Architecture** is a data design pattern that organizes data into three layers—Bronze, Silver, and Gold—representing progressive levels of data quality and refinement.

Bronze Layer (Raw Data)

Purpose: Ingest and preserve raw data in its original form

Characteristics: - Append-only, immutable data - Minimal transformation - Preserves data lineage - Enables reprocessing - Schema-on-read approach

Implementation:

```
# Bronze layer: Raw ingestion
from pyspark.sql.functions import current_timestamp, input_file_name

bronze_df = spark.readStream \
    .format("cloudFiles") \
    .option("cloudFiles.format", "json") \
    .option("cloudFiles.schemaLocation", "/mnt/schemas/events") \
    .load("/mnt/landing/events/") \
    .withColumn("ingestion_timestamp", current_timestamp()) \
    .withColumn("source_file", input_file_name())

bronze_df.writeStream \
    .format("delta") \
    .option("checkpointLocation", "/mnt/checkpoints/bronze_events") \
    .option("mergeSchema", "true") \
    .table("bronze.events")
```

Best Practices: - Store all fields as STRING or VARIANT to handle schema changes - Add metadata columns (ingestion_timestamp, source_file, source_system) - Use Auto Loader for incremental ingestion - Enable schema evolution - Retain all historical data

Silver Layer (Validated Data)

Purpose: Clean, validate, and enrich data for reliable consumption

Characteristics: - Schema enforcement - Data quality checks - Deduplication - Standardization - Type casting - Enrichment with reference data

Implementation:

```
from delta.tables import DeltaTable
from pyspark.sql.functions import *

# Silver layer: Validation and cleaning
def process_to_silver():
    # Read from bronze
    bronze_df = spark.readStream.table("bronze.events")

    # Data quality transformations
    silver_df = bronze_df \
        .filter(col("event_id").isNotNull()) \
        .filter(col("timestamp").isNotNull()) \
        .filter(col("timestamp") >= "2020-01-01") \
        .withColumn("event_date", to_date("timestamp")) \
        .withColumn("event_hour", hour("timestamp")) \
        .withColumn("user_id", col("user_id").cast("long")) \
        .withColumn("amount", col("amount").cast("decimal(10,2)")) \
        .dropDuplicates(["event_id"]) \
        .withColumn("processed_timestamp", current_timestamp())

    # Write to silver with data quality expectations
    silver_df.writeStream \
        .format("delta") \
        .outputMode("append") \
        .option("checkpointLocation", "/mnt/checkpoints/silver_events") \
        .foreachBatch(lambda batch_df, batch_id:
            write_with_quality_checks(batch_df, "silver.events")
        ) \
        .start()

def write_with_quality_checks(df, table_name):
    # Quality metrics
    total_records = df.count()
    null_user_ids = df.filter(col("user_id").isNull()).count()
    invalid_amounts = df.filter(col("amount") < 0).count()

    # Log quality metrics
    quality_metrics = spark.createDataFrame([
        {
            "table": table_name,
            "timestamp": datetime.now(),
            "total_records": total_records,
            "null_user_ids": null_user_ids,
            "invalid_amounts": invalid_amounts,
            "quality_score": 1 - ((null_user_ids + invalid_amounts) /
total_records)
        }
    ])

    quality_metrics.write.mode("append").saveAsTable("monitoring.data_quality")

    # Write validated data
    df.write.format("delta").mode("append").saveAsTable(table_name)
```

Data Quality Checks:

```
-- Add constraints to silver tables
ALTER TABLE silver.events ADD CONSTRAINT valid_amount CHECK (amount >= 0);
ALTER TABLE silver.events ADD CONSTRAINT valid_timestamp CHECK (timestamp IS NOT NULL);

-- Create expectations with Delta Live Tables
@dlt.expect_or_drop("valid_user_id", "user_id IS NOT NULL")
@dlt.expect_or_drop("valid_event_type", "event_type IN ('click', 'view', 'purchase')")
@dlt.expect_or_fail("critical_data", "amount IS NOT NULL AND amount > 0")
```

Gold Layer (Business-Level Aggregates)

Purpose: Provide curated, business-ready datasets optimized for analytics

Characteristics: - Business logic applied - Aggregated metrics - Dimensional modeling
- Optimized for query performance - Aligned with business requirements

Implementation:

```

# Gold layer: Business aggregates
from pyspark.sql.window import Window

# Customer 360 view
customer_360 = spark.sql("""
    SELECT
        c.customer_id,
        c.customer_name,
        c.customer_segment,
        c.registration_date,

        -- Transaction metrics
        COUNT(DISTINCT t.transaction_id) as total_transactions,
        SUM(t.amount) as lifetime_value,
        AVG(t.amount) as avg_transaction_value,
        MAX(t.transaction_date) as last_transaction_date,
        DATEDIFF(CURRENT_DATE(), MAX(t.transaction_date)) as
days_since_last_purchase,

        -- Product preferences
        COLLECT_LIST(DISTINCT t.product_category) as purchased_categories,

        -- Engagement metrics
        COUNT(DISTINCT e.session_id) as total_sessions,
        SUM(CASE WHEN e.event_type = 'view' THEN 1 ELSE 0 END) as total_views,
        SUM(CASE WHEN e.event_type = 'click' THEN 1 ELSE 0 END) as
total_clicks,

        -- Risk indicators
        CASE
            WHEN DATEDIFF(CURRENT_DATE(), MAX(t.transaction_date)) > 90 THEN
'High'
            WHEN DATEDIFF(CURRENT_DATE(), MAX(t.transaction_date)) > 30 THEN
'Medium'
            ELSE 'Low'
        END as churn_risk

    FROM silver.customers c
    LEFT JOIN silver.transactions t ON c.customer_id = t.customer_id
    LEFT JOIN silver.events e ON c.customer_id = e.user_id
    GROUP BY c.customer_id, c.customer_name, c.customer_segment,
c.registration_date
""")

customer_360.write.mode("overwrite").saveAsTable("gold.customer_360")

# Optimize for analytics
spark.sql("OPTIMIZE gold.customer_360 ZORDER BY (customer_id,
customer_segment)")

```

Dimensional Modeling:

```

-- Fact table: Sales transactions
CREATE TABLE gold.fact_sales (
  transaction_key BIGINT,
  date_key INT,
  customer_key INT,
  product_key INT,
  store_key INT,
  quantity INT,
  unit_price DECIMAL(10,2),
  discount_amount DECIMAL(10,2),
  tax_amount DECIMAL(10,2),
  total_amount DECIMAL(10,2)
) USING DELTA
PARTITIONED BY (date_key);

-- Dimension table: Customers
CREATE TABLE gold.dim_customer (
  customer_key INT,
  customer_id STRING,
  customer_name STRING,
  customer_segment STRING,
  customer_tier STRING,
  registration_date DATE,
  effective_date DATE,
  end_date DATE,
  is_current BOOLEAN
) USING DELTA;

-- Dimension table: Products
CREATE TABLE gold.dim_product (
  product_key INT,
  product_id STRING,
  product_name STRING,
  product_category STRING,
  product_subcategory STRING,
  brand STRING,
  supplier STRING
) USING DELTA;

-- Dimension table: Date
CREATE TABLE gold.dim_date (
  date_key INT,
  date DATE,
  day of week STRING,
  day_of_month INT,
  day of year INT,
  week_of_year INT,
  month INT,
  month name STRING,
  quarter INT,
  year INT,
  is weekend BOOLEAN,
  is_holiday BOOLEAN
) USING DELTA;

```

Performance Optimization:

```

-- Partition large tables
ALTER TABLE gold.fact_sales ADD PARTITION (date_key=20250101);

-- Z-Order for multi-dimensional clustering
OPTIMIZE gold.fact_sales ZORDER BY (customer_key, product_key);

-- Create materialized views for common queries
CREATE MATERIALIZED VIEW gold.monthly_sales_summary AS
SELECT
    d.year,
    d.month,
    d.month_name,
    p.product_category,
    c.customer_segment,
    SUM(f.total_amount) as total_revenue,
    COUNT(DISTINCT f.transaction_key) as transaction_count,
    COUNT(DISTINCT f.customer_key) as unique_customers,
    AVG(f.total_amount) as avg_transaction_value
FROM gold.fact_sales f
JOIN gold.dim_date d ON f.date_key = d.date_key
JOIN gold.dim_product p ON f.product_key = p.product_key
JOIN gold.dim_customer c ON f.customer_key = c.customer_key
GROUP BY d.year, d.month, d.month_name, p.product_category, c.customer_segment;

```

Databricks Lakebase (OLTP)

Lakebase is a fully managed PostgreSQL-based OLTP database engine integrated into the Databricks platform, enabling transactional workloads alongside analytical processing.

Architecture

Lakebase provides: - **PostgreSQL Compatibility**: Standard PostgreSQL wire protocol and SQL dialect - **Decoupled Storage**: Compute and storage separation for scalability - **Unity Catalog Integration**: Governed access to OLTP data - **Sync Tables**: Automatic synchronization to Delta Lake for analytics - **High Availability**: Built-in replication and failover

Use Cases

1. **Feature Store**: Low-latency feature serving for ML models
2. **Application State**: Store application state for Databricks Apps
3. **Real-time Data Serving**: Serve insights from gold tables to applications
4. **Operational Workflows**: Manage workflow state and orchestration metadata

Implementation Example

```
-- Create Lakebase instance
CREATE DATABASE INSTANCE ecommerce_oltp
WITH (
    instance_type = 'db.r5.xlarge',
    storage_size = 500,
    backup_retention_days = 7
);

-- Create transactional tables
CREATE TABLE ecommerce_oltp.orders (
    order_id SERIAL PRIMARY KEY,
    customer_id INTEGER NOT NULL,
    order_date TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    status VARCHAR(20) NOT NULL,
    total_amount DECIMAL(10,2) NOT NULL,
    CONSTRAINT valid_status CHECK (status IN ('pending', 'confirmed',
'shipped', 'delivered', 'cancelled'))
);

CREATE TABLE ecommerce_oltp.order_items (
    item_id SERIAL PRIMARY KEY,
    order_id INTEGER REFERENCES orders(order_id),
    product_id INTEGER NOT NULL,
    quantity INTEGER NOT NULL,
    unit_price DECIMAL(10,2) NOT NULL,
    CONSTRAINT positive_quantity CHECK (quantity > 0)
);

CREATE INDEX idx_orders_customer ON ecommerce_oltp.orders(customer_id);
CREATE INDEX idx_orders_date ON ecommerce_oltp.orders(order_date);

-- Sync to Delta Lake for analytics
CREATE SYNCED TABLE production.sales.orders_analytics
AS SELECT
    order_id,
    customer_id,
    order_date,
    status,
    total amount
FROM ecommerce_oltp.orders;

-- Synced table automatically updates as OLTP data changes
```

Integration with Databricks Apps

```
# Flask app using Lakebase
from flask import Flask, request, jsonify
import psycopg2
import os

app = Flask(__name__)

# Connect to Lakebase
def get_db_connection():
    return psycopg2.connect(
        host=os.getenv("LAKEBASE_HOST"),
        database="ecommerce_oltp",
        user=os.getenv("LAKEBASE_USER"),
        password=os.getenv("LAKEBASE_PASSWORD")
    )

@app.route('/orders', methods=['POST'])
def create_order():
    data = request.get_json()
    conn = get_db_connection()
    cursor = conn.cursor()

    try:
        # Insert order (ACID transaction)
        cursor.execute("""
            INSERT INTO orders (customer_id, status, total_amount)
            VALUES (%s, %s, %s)
            RETURNING order_id
            """, (data['customer_id'], 'pending', data['total_amount']))

        order_id = cursor.fetchone()[0]

        # Insert order items
        for item in data['items']:
            cursor.execute("""
                INSERT INTO order_items (order_id, product_id, quantity,
unit_price)
                VALUES (%s, %s, %s, %s)
                """, (order_id, item['product_id'], item['quantity'],
item['unit_price']))

        conn.commit()
        return jsonify({"order_id": order_id, "status": "created"}), 201

    except Exception as e:
        conn.rollback()
        return jsonify({"error": str(e)}), 500
    finally:
        cursor.close()
        conn.close()

@app.route('/orders/<int:order_id>', methods=['GET'])
def get_order(order_id):
    conn = get_db_connection()
    cursor = conn.cursor()

    cursor.execute("""
        SELECT o.order_id, o.customer_id, o.order_date, o.status,
o.total_amount,
        json_agg(
            json build object(
                'product_id', oi.product_id,
                'quantity', oi.quantity,
                'unit_price', oi.unit_price
            )
        ) as items
        FROM orders o
        LEFT JOIN order_items oi ON o.order_id = oi.order_id
        WHERE o.order_id = %s
        """, (order_id,))

    rows = cursor.fetchall()
    if not rows:
        return jsonify({"error": "Order not found"}), 404

    order = {
        "order_id": rows[0][0],
        "customer_id": rows[0][1],
        "order_date": rows[0][2],
        "status": rows[0][3],
        "total_amount": rows[0][4],
        "items": rows[0][5]
    }

    return jsonify(order), 200
```

```

        )) as items
    FROM orders o
    LEFT JOIN order_items oi ON o.order_id = oi.order_id
    WHERE o.order_id = %s
    GROUP BY o.order_id
    """, (order_id,))

result = cursor.fetchone()
cursor.close()
conn.close()

if result:
    return jsonify({
        "order_id": result[0],
        "customer_id": result[1],
        "order_date": result[2].isoformat(),
        "status": result[3],
        "total_amount": float(result[4]),
        "items": result[5]
    }), 200
else:
    return jsonify({"error": "Order not found"}), 404

```

Implementation Guide

Step 1: Set Up Unity Catalog

```

-- Create catalog hierarchy
CREATE CATALOG IF NOT EXISTS production;
CREATE SCHEMA IF NOT EXISTS production.bronze;
CREATE SCHEMA IF NOT EXISTS production.silver;
CREATE SCHEMA IF NOT EXISTS production.gold;

-- Set up permissions
GRANT USE CATALOG ON CATALOG production TO `data_team`;
GRANT CREATE SCHEMA ON CATALOG production TO `data_engineers`;
GRANT SELECT ON SCHEMA production.gold TO `data_analysts`;

```


Step 2: Implement Bronze Layer

```
# Auto Loader for continuous ingestion
from pyspark.sql.functions import *

# Configure Auto Loader
bronze_stream = spark.readStream \
    .format("cloudFiles") \
    .option("cloudFiles.format", "json") \
    .option("cloudFiles.schemaLocation", "/mnt/schemas/events") \
    .option("cloudFiles.inferColumnTypes", "true") \
    .option("cloudFiles.schemaEvolutionMode", "addNewColumns") \
    .load("s3://my-bucket/landing/events/")

# Add metadata
bronze_enriched = bronze_stream \
    .withColumn("_ingestion_timestamp", current_timestamp()) \
    .withColumn("_source_file", input_file_name()) \
    .withColumn("_bronze_date", current_date())

# Write to bronze
bronze_enriched.writeStream \
    .format("delta") \
    .option("checkpointLocation", "/mnt/checkpoints/bronze_events") \
    .option("mergeSchema", "true") \
    .partitionBy("_bronze_date") \
    .table("production.bronze.events")
```

Step 3: Implement Silver Layer with DLT

```
import dlt
from pyspark.sql.functions import *

@dlt.table(
    name="silver_events",
    comment="Validated and cleaned events",
    table_properties={"quality": "silver", "pipelines.autoOptimize.zOrderCols":
"user_id,event_date"}
)
@dlt.expect_or_drop("valid_event_id", "event_id IS NOT NULL")
@dlt.expect_or_drop("valid_timestamp", "event_timestamp IS NOT NULL AND
event_timestamp >= '2020-01-01'")
@dlt.expect_or_drop("valid_user", "user_id IS NOT NULL AND user_id > 0")
@dlt.expect("valid_amount", "amount >= 0")
def create_silver_events():
    return (
        dlt.read_stream("production.bronze.events")
        .select(
            col("event_id"),
            col("user_id").cast("long"),
            to_timestamp("event_timestamp").alias("event_timestamp"),
            to_date("event_timestamp").alias("event_date"),
            col("event_type"),
            col("amount").cast("decimal(10,2)"),
            col("_ingestion_timestamp")
        )
        .dropDuplicates(["event_id"])
        .withColumn("_silver_processed_at", current_timestamp())
    )
```

Step 4: Implement Gold Layer

```
@dlt.table(
    name="gold_daily_user_metrics",
    comment="Daily aggregated user metrics for analytics",
    table_properties={"quality": "gold"}
)
def create_gold_daily_metrics():
    return (
        dlt.read("silver_events")
        .groupBy("user_id", "event_date")
        .agg(
            count("*").alias("total_events"),
            countDistinct("event_id").alias("unique_events"),
            sum(when(col("event_type") == "purchase",
col("amount")).otherwise(0)).alias("total_revenue"),
            count(when(col("event_type") == "purchase",
1)).alias("purchase_count"),
            count(when(col("event_type") == "view", 1)).alias("view_count"),
            count(when(col("event_type") == "click", 1)).alias("click_count")
        )
        .withColumn("conversion_rate",
            col("purchase_count") / (col("view_count") + col("click_count")))
    )
```

Step 5: Optimize and Monitor

```
-- Optimize tables regularly
OPTIMIZE production.silver.events ZORDER BY (user_id, event_date);
OPTIMIZE production.gold.daily_user_metrics ZORDER BY (event_date, user_id);

-- Vacuum old files (retain 7 days)
VACUUM production.silver.events RETAIN 168 HOURS;

-- Monitor data quality
SELECT
    table_name,
    COUNT(*) as row_count,
    COUNT(DISTINCT user_id) as unique_users,
    MIN(event_date) as earliest_date,
    MAX(event_date) as latest_date
FROM production.silver.events
GROUP BY table_name;

-- Set up table monitoring
CREATE OR REPLACE TABLE production.monitoring.table_metrics AS
SELECT
    current_timestamp() as check_timestamp,
    'production.silver.events' as table_name,
    COUNT(*) as row_count,
    COUNT(DISTINCT user_id) as unique_users,
    SUM(CASE WHEN amount IS NULL THEN 1 ELSE 0 END) as null_amounts
FROM production.silver.events;
```

Databricks Asset Bundles

Overview and Concepts

Databricks Asset Bundles (DAB) provide an infrastructure-as-code approach to managing Databricks resources. Bundles enable developers to define, version, validate, and deploy Databricks workflows, apps, pipelines, and other resources programmatically.

Key Benefits

1. **Infrastructure as Code:** Define all Databricks resources in YAML or Python
2. **Version Control:** Track changes in Git alongside application code
3. **CI/CD Integration:** Automate testing and deployment pipelines
4. **Environment Management:** Separate dev, staging, and production configurations
5. **Reproducibility:** Ensure consistent deployments across environments
6. **Collaboration:** Enable team-based development with code reviews

Core Concepts

- **Bundle:** Collection of Databricks resources and their configurations
- **Resources:** Jobs, pipelines, apps, models, dashboards, etc.
- **Targets:** Environment-specific configurations (dev, staging, prod)
- **Variables:** Parameterized values for flexibility
- **Deployment Modes:** Controls resource naming and permissions

Bundle Structure

A typical bundle project structure:

```
my_project/
├── databricks.yml          # Main bundle configuration
├── resources/              # Resource definitions
│   ├── jobs/
│   │   ├── etl_job.yml
│   │   └── ml_training_job.yml
│   ├── pipelines/
│   │   └── data_pipeline.yml
│   └── apps/
│       └── analytics_app.yml
├── src/                    # Source code
│   ├── notebooks/
│   │   ├── bronze_ingestion.py
│   │   ├── silver_transformation.py
│   │   └── gold_aggregation.py
│   └── python/
│       ├── __init__.py
│       └── utils.py
├── tests/                  # Unit tests
│   └── test_transformations.py
├── fixtures/              # Test data
│   └── sample_data.json
└── README.md
```

Configuration Reference

databricks.yml Structure

```
# Bundle definition
bundle:
  name: my_data_platform

# Git integration
git:
  origin_url: https://github.com/myorg/my-data-platform
  branch: main

# Variables for parameterization
variables:
  catalog_name:
    description: "Unity Catalog name"
    default: "development"

  warehouse_id:
    description: "SQL Warehouse ID"
    default: "abc123def456"

  notification_email:
    description: "Email for job notifications"

# Include additional configuration files
include:
  - resources/**/* .yml

# Workspace configuration
workspace:
  host: https://company.databricks.com
  root_path:
    /Workspace/Users/${workspace.current_user.userName}/.bundle/${bundle.name}/${bu

# Define resources
resources:
  jobs:
    etl_pipeline:
      name: "[${bundle.target}] ETL Pipeline"
      tasks:
        - task_key: bronze_ingestion
          notebook_task:
            notebook_path: ../src/notebooks/bronze_ingestion.py
            source: WORKSPACE
          new_cluster:
            spark_version: "13.3.x-scala2.12"
            node_type_id: "i3.xlarge"
            num_workers: 2

        - task_key: silver_transformation
          depends_on:
            - task_key: bronze_ingestion
          notebook_task:
            notebook_path: ../src/notebooks/silver_transformation.py
          new_cluster:
            spark_version: "13.3.x-scala2.12"
            node_type_id: "i3.xlarge"
            num_workers: 4

  schedule:
    quartz_cron_expression: "0 0 2 * * ?"
```

```

    timezone_id: "America/Los_Angeles"

    email_notifications:
      on_failure:
        - ${var.notification_email}

pipelines:
  dlt_pipeline:
    name: "[${bundle.target}] DLT Pipeline"
    catalog: ${var.catalog_name}
    target: ${bundle.target}_schema
    libraries:
      - notebook:
          path: ../src/notebooks/dlt_definitions.py
    configuration:
      warehouse_id: ${var.warehouse_id}
    continuous: false

apps:
  analytics_dashboard:
    name: "[${bundle.target}] Analytics Dashboard"
    description: "Real-time analytics dashboard"
    resources:
      - name: warehouse
        sql_warehouse:
          id: ${var.warehouse_id}

# Deployment targets
targets:
  dev:
    mode: development
    default: true
    workspace:
      host: https://dev.databricks.com
    variables:
      catalog_name: "dev"
      notification_email: "dev-team@company.com"

  staging:
    mode: development
    workspace:
      host: https://staging.databricks.com
    variables:
      catalog_name: "staging"
      notification_email: "qa-team@company.com"

  prod:
    mode: production
    workspace:
      host: https://prod.databricks.com
    run_as:
      service_principal_name: "prod-service-principal"
    variables:
      catalog_name: "production"
      notification_email: "data-ops@company.com"

# Production-specific overrides
resources:
  jobs:
    etl_pipeline:
      schedule:
        quartz_cron_expression: "0 0 1 * * ?" # Run at 1 AM in prod
      tasks:
        - task_key: bronze_ingestion
          new_cluster:
            num_workers: 8 # More workers in prod

```

```
- task_key: silver_transformation
  new_cluster:
    num_workers: 16
```

Development Lifecycle

1. Initialize Bundle

```
# Initialize new bundle from template
databricks bundle init

# Choose template
# - default-python: Python-based workflows
# - default-sql: SQL-based workflows
# - dbt-sql: dbt integration
# - mlops-stacks: ML workflows

# Or initialize from custom template
databricks bundle init --template-dir /path/to/template
```

2. Develop Locally

```
# Validate bundle configuration
databricks bundle validate

# Preview deployment changes
databricks bundle deploy --dry-run

# Deploy to development
databricks bundle deploy --target dev
```

3. Test and Iterate

```
# Run specific job
databricks bundle run etl_pipeline --target dev

# Run with parameters
databricks bundle run etl_pipeline \
  --target dev \
  --params '{"start_date": "2025-01-01", "end_date": "2025-01-31"}'

# View job status
databricks bundle run etl_pipeline --target dev --wait
```

4. Validate and Deploy

```
# Run validation
databricks bundle validate --target prod

# Deploy to production
databricks bundle deploy --target prod

# Verify deployment
databricks bundle summary --target prod
```

5. Clean Up

```
# Destroy resources (use with caution!)
databricks bundle destroy --target dev

# Confirm destruction
# Type 'yes' when prompted
```


CI/CD Integration

GitHub Actions Workflow

```
# .github/workflows/databricks-deploy.yml
name: Databricks Bundle Deployment

on:
  push:
    branches: [main, develop]
  pull_request:
    branches: [main]

env:
  DATABRICKS_HOST: ${ secrets.DATABRICKS_HOST }
  DATABRICKS_TOKEN: ${ secrets.DATABRICKS_TOKEN }

jobs:
  validate:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3

      - name: Install Databricks CLI
        run: |
          curl -fsSL https://raw.githubusercontent.com/databricks/setup-
cli/main/install.sh | sh

      - name: Validate Bundle
        run: databricks bundle validate

  test:
    runs-on: ubuntu-latest
    needs: validate
    steps:
      - uses: actions/checkout@v3

      - name: Set up Python
        uses: actions/setup-python@v4
        with:
          python-version: '3.10'

      - name: Install dependencies
        run: |
          pip install -r requirements.txt
          pip install pytest

      - name: Run tests
        run: pytest tests/

  deploy-dev:
    runs-on: ubuntu-latest
    needs: [validate, test]
    if: github.ref == 'refs/heads/develop'
    steps:
      - uses: actions/checkout@v3

      - name: Install Databricks CLI
        run: |
          curl -fsSL https://raw.githubusercontent.com/databricks/setup-
cli/main/install.sh | sh

      - name: Deploy to Dev
```

```
    run: |
      databricks bundle deploy --target dev
      databricks bundle run etl_pipeline --target dev

deploy-prod:
  runs-on: ubuntu-latest
  needs: [validate, test]
  if: github.ref == 'refs/heads/main'
  environment: production
  steps:
    - uses: actions/checkout@v3

    - name: Install Databricks CLI
      run: |
        curl -fsSL https://raw.githubusercontent.com/databricks/setup-
        cli/main/install.sh | sh

    - name: Deploy to Production
      run: databricks bundle deploy --target prod
```

Azure DevOps Pipeline

```
# azure-pipelines.yml
trigger:
  branches:
    include:
      - main
      - develop

pool:
  vmImage: 'ubuntu-latest'

variables:
  - group: databricks-credentials

stages:
  - stage: Validate
    jobs:
      - job: ValidateBundle
        steps:
          - task: UsePythonVersion@0
            inputs:
              versionSpec: '3.10'

          - script: |
              curl -fsSL https://raw.githubusercontent.com/databricks/setup-
cli/main/install.sh | sh
              databricks bundle validate
            displayName: 'Validate Databricks Bundle'
            env:
              DATABRICKS_HOST: $(DATABRICKS_HOST)
              DATABRICKS_TOKEN: $(DATABRICKS_TOKEN)

  - stage: Test
    dependsOn: Validate
    jobs:
      - job: RunTests
        steps:
          - task: UsePythonVersion@0
            inputs:
              versionSpec: '3.10'

          - script: |
              pip install -r requirements.txt
              pip install pytest
              pytest tests/
            displayName: 'Run Unit Tests'

  - stage: DeployDev
    dependsOn: Test
    condition: eq(variables['Build.SourceBranch'], 'refs/heads/develop')
    jobs:
      - deployment: DeployToDev
        environment: 'development'
        strategy:
          runOnce:
            deploy:
              steps:
                - checkout: self
                - script: |
                    curl -fsSL
https://raw.githubusercontent.com/databricks/setup-cli/main/install.sh | sh
                    databricks bundle deploy --target dev
                  displayName: 'Deploy to Development'
                  env:
```

```

        DATABRICKS_HOST: $(DATABRICKS_DEV_HOST)
        DATABRICKS_TOKEN: $(DATABRICKS_DEV_TOKEN)

- stage: DeployProd
  dependsOn: Test
  condition: eq(variables['Build.SourceBranch'], 'refs/heads/main')
  jobs:
    - deployment: DeployToProduction
      environment: 'production'
      strategy:
        runOnce:
          deploy:
            steps:
              - checkout: self
              - script: |
                curl -fsSL
                https://raw.githubusercontent.com/databricks/setup-cli/main/install.sh | sh
                databricks bundle deploy --target prod
              displayName: 'Deploy to Production'
            env:
              DATABRICKS_HOST: $(DATABRICKS_PROD_HOST)
              DATABRICKS_TOKEN: $(DATABRICKS_PROD_TOKEN)

```

Deployment Modes

Databricks Asset Bundles support three deployment modes:

1. Development Mode

```

targets:
  dev:
    mode: development

```

Characteristics: - Resources prefixed with [dev username] - Job schedules paused by default - Permissions: Only creator can access - Use case: Individual developer environments

2. Production Mode

```

targets:
  prod:
    mode: production
    run_as:
      service_principal_name: "prod-sp"

```

Characteristics: - No resource name prefixes - Job schedules active - Runs as service principal - Immutable deployments (prevents accidental changes) - Use case: Production workloads

3. Snapshot Mode (Default)

```
targets:
  staging:
    mode: snapshot
```

Characteristics: - Resources prefixed with target name - Job schedules active - Runs as deploying user - Use case: Staging/QA environments

Best Practices

1. Project Organization

```
my_project/
├── databricks.yml           # Main config
├── resources/
│   ├── common.yml          # Shared configurations
│   ├── jobs/
│   │   ├── ingestion.yml
│   │   └── transformation.yml
│   └── pipelines/
│       └── dlt_pipeline.yml
├── src/
│   ├── common/              # Shared utilities
│   ├── ingestion/           # Ingestion logic
│   └── transformation/      # Transformation logic
├── tests/
│   ├── unit/
│   └── integration/
├── docs/
│   └── architecture.md
```

2. Use Variables for Flexibility

```
variables:
  environment:
    description: "Deployment environment"

  cluster_config:
    description: "Cluster configuration"
    default:
      spark_version: "13.3.x-scala2.12"
      node_type_id: "i3.xlarge"

resources:
  jobs:
    my_job:
      name: "[${var.environment}] My Job"
      new_cluster: ${var.cluster_config}
```

3. Separate Concerns

```
# resources/jobs/etl.yml
resources:
  jobs:
    etl_job:
      name: "ETL Job"
      tasks: !include tasks/etl_tasks.yml

# resources/tasks/etl_tasks.yml
- task_key: ingest
  notebook_task:
    notebook_path: ../src/ingest.py
- task_key: transform
  depends_on:
    - task_key: ingest
  notebook_task:
    notebook_path: ../src/transform.py
```

4. Version Control Best Practices

- Store bundles in Git
- Use feature branches for development
- Require code reviews for production changes
- Tag releases for production deployments
- Use `.gitignore` for generated files

```
# .gitignore
.databricks/
__pycache__/
*.pyc
.pytest_cache/
.venv/
```

5. Testing Strategy

```
# tests/test_transformations.py
import pytest
from pyspark.sql import SparkSession
from src.transformations import clean_data

@pytest.fixture
def spark():
    return SparkSession.builder.master("local[1]").getOrCreate()

def test_clean_data(spark):
    # Arrange
    input_data = [
        (1, "John", None),
        (2, "Jane", "invalid"),
        (3, "Bob", "valid")
    ]
    df = spark.createDataFrame(input_data, ["id", "name", "status"])

    # Act
    result = clean_data(df)

    # Assert
    assert result.count() == 2 # Invalid records dropped
    assert result.filter("status = 'valid'").count() == 1
```

Code Examples

Example 1: Complete ETL Bundle

databricks.yml:

```
bundle:
  name: customer_analytics_etl

variables:
  catalog:
    default: "development"
  schema:
    default: "customer_data"

include:
  - resources/*.yaml

targets:
  dev:
    mode: development
    default: true
    variables:
      catalog: "dev"

  prod:
    mode: production
    run as:
      service_principal_name: "etl-service-principal"
    variables:
      catalog: "prod"
```

resources/jobs.yml:


```

resources:
  jobs:
    customer_etl:
      name: "[${bundle.target}] Customer Analytics ETL"

      job_clusters:
        - job_cluster_key: "etl_cluster"
          new_cluster:
            spark_version: "13.3.x-scala2.12"
            node_type_id: "i3.xlarge"
            num_workers: 4
            spark_conf:
              "spark.databricks.delta.optimizeWrite.enabled": "true"
              "spark.databricks.delta.autoCompact.enabled": "true"

      tasks:
        - task_key: "ingest_raw_data"
          job_cluster_key: "etl_cluster"
          notebook_task:
            notebook_path: "../src/notebooks/01_ingest_raw.py"
            base_parameters:
              catalog: "${var.catalog}"
              schema: "${var.schema}"

        - task_key: "transform_to_silver"
          depends_on:
            - task_key: "ingest_raw_data"
          job_cluster_key: "etl_cluster"
          notebook_task:
            notebook_path: "../src/notebooks/02_transform_silver.py"
            base_parameters:
              catalog: "${var.catalog}"
              schema: "${var.schema}"

        - task_key: "aggregate_to_gold"
          depends_on:
            - task_key: "transform_to_silver"
          job_cluster_key: "etl_cluster"
          notebook_task:
            notebook_path: "../src/notebooks/03_aggregate_gold.py"
            base_parameters:
              catalog: "${var.catalog}"
              schema: "${var.schema}"

        - task_key: "data_quality_checks"
          depends_on:
            - task_key: "aggregate_to_gold"
          job_cluster_key: "etl_cluster"
          python_wheel_task:
            package_name: "data_quality"
            entry_point: "run_checks"
            parameters:
              - "--catalog=${var.catalog}"
              - "--schema=${var.schema}"
          libraries:
            - whl: "../dist/data_quality-0.1.0-py3-none-any.whl"

      schedule:
        quartz_cron_expression: "0 0 2 * * ?"
        timezone_id: "UTC"
        pause_status: "UNPAUSED"

      email_notifications:
        on_failure:
          - "data-eng@company.com"

```

```
on_success:  
  - "data-ops@company.com"
```

```
max_concurrent_runs: 1  
timeout_seconds: 7200
```

Example 2: ML Training Pipeline Bundle

```
bundle:
  name: ml_training_pipeline

variables:
  model_name:
    default: "customer_churn_model"
  experiment_path:
    default: "/Shared/ml_experiments/customer_churn"

resources:
  jobs:
    train_model:
      name: "[$`${bundle.target}`] ML Training - `${var.model_name}`"

      tasks:
        - task_key: "prepare_features"
          new_cluster:
            spark_version: "13.3.x-cpu-ml-scala2.12"
            node_type_id: "i3.xlarge"
            num_workers: 2
          notebook_task:
            notebook_path: "../src/ml/01_feature_engineering.py"
            base_parameters:
              model_name: "${var.model_name}"

        - task_key: "train_model"
          depends_on:
            - task_key: "prepare_features"
          new_cluster:
            spark_version: "13.3.x-cpu-ml-scala2.12"
            node_type_id: "i3.2xlarge"
            num_workers: 4
          notebook_task:
            notebook_path: "../src/ml/02_train_model.py"
            base_parameters:
              model_name: "${var.model_name}"
              experiment_path: "${var.experiment_path}"
          libraries:
            - pypi:
                package: "scikit-learn==1.3.0"
            - pypi:
                package: "xgboost==2.0.0"

        - task_key: "evaluate_model"
          depends_on:
            - task_key: "train_model"
          new_cluster:
            spark_version: "13.3.x-cpu-ml-scala2.12"
            node_type_id: "i3.xlarge"
            num_workers: 2
          notebook_task:
            notebook_path: "../src/ml/03_evaluate_model.py"
            base_parameters:
              model_name: "${var.model_name}"
              experiment_path: "${var.experiment_path}"

        - task_key: "register_model"
          depends_on:
            - task_key: "evaluate_model"
          new_cluster:
            spark_version: "13.3.x-cpu-ml-scala2.12"
            node_type_id: "i3.xlarge"
            num_workers: 1
```

```

    notebook_task:
      notebook_path: "../src/ml/04_register_model.py"
      base_parameters:
        model_name: "${var.model_name}"
        experiment_path: "${var.experiment_path}"
        registry_stage: "${bundle.target == 'prod' ? 'Production' :
'Staging'}"

    experiments:
      churn_experiment:
        name: "${var.experiment_path}"
        description: "Customer churn prediction experiments"

    models:
      churn_model:
        name: "${var.model_name}"
        description: "Customer churn prediction model"

  targets:
    dev:
      mode: development
      variables:
        model_name: "dev_customer_churn"
        experiment_path:
"/Users/${workspace.current_user.userName}/experiments/churn"

    prod:
      mode: production
      run_as:
        service_principal_name: "ml-service-principal"
      variables:
        model_name: "customer_churn_model"
        experiment_path: "/Shared/ml_experiments/customer_churn"

```

Example 3: Multi-App Bundle

```
bundle:
  name: analytics_platform

resources:
  apps:
    sales_dashboard:
      name: "[${bundle.target}] Sales Dashboard"
      description: "Real-time sales analytics"
      resources:
        - name: sales_warehouse
          sql_warehouse:
            id: "${var.warehouse_id}"

    customer_insights:
      name: "[${bundle.target}] Customer Insights"
      description: "Customer behavior analytics"
      resources:
        - name: analytics_warehouse
          sql_warehouse:
            id: "${var.warehouse_id}"

    ml_predictions:
      name: "[${bundle.target}] ML Predictions"
      description: "Real-time ML predictions interface"
      resources:
        - name: model_endpoint
          model_serving_endpoint:
            name: "customer-churn-endpoint"

  jobs:
    refresh_dashboards:
      name: "[${bundle.target}] Refresh Dashboard Data"
      tasks:
        - task_key: "refresh_sales"
          sql_task:
            warehouse_id: "${var.warehouse_id}"
            query:
              query_id: "${var.sales_query_id}"

        - task_key: "refresh_customers"
          sql_task:
            warehouse_id: "${var.warehouse_id}"
            query:
              query_id: "${var.customer_query_id}"

      schedule:
        quartz_cron_expression: "0 */15 * * * ?" # Every 15 minutes
        timezone_id: "UTC"

  variables:
    warehouse_id:
      description: "SQL Warehouse ID for apps"
    sales_query_id:
      description: "Sales refresh query ID"
    customer_query_id:
      description: "Customer refresh query ID"

  targets:
    dev:
      mode: development
      variables:
        warehouse_id: "dev_warehouse_123"
        sales_query_id: "dev_sales_query"
```

```

        customer_query_id: "dev_customer_query"

prod:
  mode: production
  run_as:
    service_principal_name: "apps-service-principal"
  variables:
    warehouse_id: "prod_warehouse_456"
    sales_query_id: "prod_sales_query"
    customer_query_id: "prod_customer_query"

```

Integration Patterns

Pattern 1: Apps + Lakehouse

Databricks Apps can leverage the full power of the Lakehouse architecture:

```

# Streamlit app accessing Lakehouse data
import streamlit as st
from databricks import sql
import os

# Connect to SQL Warehouse
connection = sql.connect(
    server_hostname=os.getenv("DATABRICKS_HOST"),
    http_path=f"/sql/1.0/warehouses/{os.getenv('WAREHOUSE_ID')}"
)

# Query gold layer
@st.cache_data(ttl=600)
def load_customer_360(customer_id):
    cursor = connection.cursor()
    cursor.execute(f"""
        SELECT *
        FROM production.gold.customer_360
        WHERE customer_id = '{customer_id}'
    """)
    return cursor.fetchall_arrow().to_pandas()

# UI
customer_id = st.text_input("Customer ID")
if customer_id:
    df = load_customer_360(customer_id)
    st.dataframe(df)

```

Pattern 2: Bundles + Apps

Manage Databricks Apps as code using Asset Bundles:

```
# databricks.yml
resources:
  apps:
    analytics_app:
      name: "[${bundle.target}] Analytics App"
      description: "Customer analytics dashboard"
      resources:
        - name: warehouse
          sql_warehouse:
            id: "${var.warehouse_id}"
        - name: lakebase
          database_instance:
            name: "analytics_db"

# Deploy app with bundle
# databricks bundle deploy --target prod
```

Pattern 3: Bundles + Lakehouse

Orchestrate complete Lakehouse workflows with bundles:

```
resources:
  pipelines:
    medallion_pipeline:
      name: "Medallion Architecture Pipeline"
      catalog: "${var.catalog}"
      target: "${bundle.target}_schema"
      libraries:
        - notebook:
            path: "../src/bronze_layer.py"
        - notebook:
            path: "../src/silver_layer.py"
        - notebook:
            path: "../src/gold_layer.py"

      configuration:
        bronze_path: "/mnt/bronze"
        silver_path: "/mnt/silver"
        gold_path: "/mnt/gold"
```

Pattern 4: End-to-End Platform

Complete integration of all three technologies:

```

bundle:
  name: customer_intelligence_platform

resources:
  # Data ingestion and transformation
  pipelines:
    customer_data_pipeline:
      name: "Customer Data Pipeline"
      catalog: "production"
      target: "customer_data"
      libraries:
        - notebook:
            path: "../src/pipelines/bronze_customers.py"
        - notebook:
            path: "../src/pipelines/silver_customers.py"
        - notebook:
            path: "../src/pipelines/gold_customer_360.py"

  # ML training
  jobs:
    churn_model_training:
      name: "Churn Model Training"
      tasks:
        - task_key: "train"
          notebook_task:
            notebook_path: "../src/ml/train_churn_model.py"

  # Model serving
  model_serving_endpoints:
    churn_prediction:
      name: "churn-prediction-endpoint"
      config:
        served_entities:
          - entity_name: "production.ml_models.customer_churn"
            entity_version: "1"
            workload_size: "Small"

  # OLTP database for app state
  database_instances:
    app_state_db:
      name: "customer app state"
      instance_type: "db.t3.medium"

  # Customer-facing application
  apps:
    customer_portal:
      name: "Customer Intelligence Portal"
      description: "360-degree customer view with churn predictions"
      resources:
        - name: warehouse
          sql_warehouse:
            id: "${var.warehouse_id}"
        - name: model
          model_serving_endpoint:
            name: "churn-prediction-endpoint"
        - name: database
          database_instance:
            name: "customer_app_state"

```


References

Official Databricks Documentation

1. **Databricks Apps**
2. [Databricks Apps Overview](#)
3. [Get Started with Databricks Apps](#)
4. [Develop Databricks Apps](#)
5. [Configure App Runtime](#)
6. **Databricks Lakehouse**
7. [Lakehouse Architecture](#)
8. [Lakehouse Reference Architecture](#)
9. [Medallion Architecture](#)
10. [What is Lakebase?](#)
11. **Databricks Asset Bundles**
12. [What are Databricks Asset Bundles?](#)
13. [Bundle Configuration](#)
14. [Develop Bundles](#)
15. [Bundle Tutorials](#)
16. [CI/CD Best Practices](#)

GitHub Repositories

1. **Databricks Apps Examples**
2. [databricks-solutions/databricks-apps-examples](#)
3. **Databricks Asset Bundle Examples**
4. [databricks/bundle-examples](#)

Community Resources

1. **Medium Articles**
2. [Building a Databricks App Project: Architecture, Concepts, and Implementation Guide](#)
3. [CI/CD Strategies For Databricks Asset Bundles](#)
4. [Understanding Databricks Lakehouse Reference Architectures](#)
5. **Databricks Community**
6. [Exploring Code With Databricks Apps](#)

Additional Resources

- [Databricks Glossary: Medallion Architecture](#)
 - [Databricks Product Page: Databricks Apps](#)
 - [Databricks Blog: Announcing General Availability of Databricks Asset Bundles](#)
-

Document Version: 1.0

Last Updated: October 20, 2025

Compiled By: Comprehensive Research Analysis

This document represents a complete synthesis of official Databricks documentation, community resources, and open-source examples to provide actionable guidance for implementing Databricks Apps, Lakehouse architecture, and Asset Bundles in production environments.

Databricks Unity Catalog

Overview and Architecture

Databricks Unity Catalog is a unified governance solution for data and AI assets on the Databricks platform. It provides centralized access control, auditing, lineage tracking, quality monitoring, and data discovery capabilities across all Databricks workspaces in an organization.

Unity Catalog represents a paradigm shift in data governance by offering a single, consistent governance layer that spans across clouds, regions, and workspaces. It eliminates the complexity of managing multiple governance systems and provides a standards-compliant security model based on ANSI SQL.

Core Value Proposition:

Unity Catalog addresses the fundamental challenge of modern data governance: how to maintain security, compliance, and discoverability while enabling data democratization at scale. It provides enterprise-grade governance without sacrificing agility or developer productivity.

Key Features

1. Define Once, Secure Everywhere

Unity Catalog offers a single place to administer data access policies that apply across all workspaces in a region. This eliminates the need to replicate permissions across multiple systems and ensures consistent security posture.

2. Standards-Compliant Security Model

The security model is based on standard ANSI SQL, allowing administrators to grant permissions using familiar syntax. This reduces the learning curve and enables integration with existing governance frameworks.

3. Built-in Auditing and Lineage

Unity Catalog automatically captures user-level audit logs that record all access to data assets. It also captures comprehensive lineage data that tracks how data assets are created, transformed, and used across all languages (SQL, Python, R, Scala).

4. Data Discovery

Unity Catalog provides tagging, documentation, and search capabilities to help data consumers find the data they need. This includes support for business glossaries, data quality metrics, and usage statistics.

5. System Tables

Unity Catalog provides access to operational data through system tables, including audit logs, billable usage, and lineage information. This enables organizations to build

custom monitoring and governance solutions.

Architecture Components

The Three-Level Namespace

Unity Catalog uses a three-level hierarchy for organizing data and AI assets:

Level 1: Metastore - Top-level container for metadata - One metastore per region recommended - Multi-tenant service boundary - Registers metadata about data and AI assets

Level 2: Catalogs - Organize data assets by business domain or lifecycle - Typically mirror organizational units (e.g., marketing, finance, engineering) - Can represent environments (e.g., dev, staging, production)

Level 3: Schemas (Databases) - Contain tables, views, volumes, models, and functions - Organize assets into logical categories - Typically represent projects, use cases, or team sandboxes

Level 4: Data and AI Objects - **Tables:** Managed or external collections of structured data - **Views:** Saved queries against tables - **Volumes:** Storage for unstructured data (files, images, documents) - **Functions:** User-defined functions (UDFs) - **Models:** ML models registered with MLflow

Securable Objects

Unity Catalog manages access through several types of securable objects:

Data Access Objects: - **Storage Credentials:** Encapsulate cloud credentials for accessing storage - **External Locations:** Reference cloud storage paths with associated credentials - **Connections:** Provide access to external databases via Lakehouse Federation - **Service Credentials:** Access external services

Sharing Objects: - **Shares:** Collections of data shared via Delta Sharing - **Recipients:** Entities that receive shared data - **Providers:** Entities that share data - **Clean Rooms:** Secure collaboration environments

Access Control Model

Admin Roles

Unity Catalog defines three primary admin roles:

Role	Scope	Key Privileges
Account Admin	Account-wide	Create metastores, link workspaces, manage users
Metastore Admin	Metastore	Manage storage, create catalogs, grant privileges
Workspace Admin	Workspace	Manage workspace objects, add users, configure compute

Privilege Hierarchy

Privileges in Unity Catalog follow an inheritance model:

```
Metastore
├─ Catalog (USAGE, CREATE SCHEMA, MANAGE)
│   └─ Schema (USAGE, CREATE TABLE, CREATE FUNCTION)
│       └─ Table/View (SELECT, MODIFY, READ_METADATA)
```

Access to a parent object implicitly grants the same access to all children unless explicitly revoked.

Grant Syntax

Unity Catalog uses standard SQL GRANT and REVOKE statements:

```
-- Grant table access to a group
GRANT SELECT ON TABLE catalog.schema.table TO `data-analysts`;

-- Grant schema creation privileges
GRANT CREATE SCHEMA ON CATALOG production TO `data-engineers`;

-- Grant catalog usage to all users
GRANT USAGE ON CATALOG shared_data TO `account users`;

-- Revoke privileges
REVOKE SELECT ON TABLE catalog.schema.sensitive_data FROM `contractors`;
```

Implementation Guide

Step 1: Metastore Setup

```
-- Create a metastore (Account Admin only)
-- This is typically done through the Databricks UI or Terraform

-- Verify metastore attachment
SELECT CURRENT_METASTORE();

-- Check metastore details
DESCRIBE METASTORE;
```

Step 2: Create Catalog Structure

```
-- Create catalogs for different environments
CREATE CATALOG IF NOT EXISTS development
  COMMENT 'Development environment catalog';

CREATE CATALOG IF NOT EXISTS production
  COMMENT 'Production environment catalog';

-- Create schemas within catalogs
CREATE SCHEMA IF NOT EXISTS production.customer_data
  COMMENT 'Customer data and analytics';

CREATE SCHEMA IF NOT EXISTS production.ml_models
  COMMENT 'Production ML models';
```

Step 3: Configure Storage

```
-- Create storage credential
CREATE STORAGE CREDENTIAL aws_s3_credential
  WITH (
    AWS_IAM_ROLE = 'arn:aws:iam::123456789012:role/databricks-s3-access'
  )
  COMMENT 'S3 access for production data';

-- Create external location
CREATE EXTERNAL LOCATION production_data
  URL 's3://my-company-data/production/'
  WITH (STORAGE CREDENTIAL aws s3 credential)
  COMMENT 'Production data storage location';

-- Set managed storage location for catalog
ALTER CATALOG production
  SET MANAGED LOCATION 's3://my-company-data/managed/production/';
```

Step 4: Create and Manage Tables

```
-- Create managed table (Unity Catalog manages lifecycle)
CREATE TABLE production.customer_data.customers (
  customer_id BIGINT,
  name STRING,
  email STRING,
  created_at TIMESTAMP
)
USING DELTA
COMMENT 'Customer master data';

-- Create external table (data managed externally)
CREATE EXTERNAL TABLE production.customer_data.transactions
LOCATION 's3://my-company-data/transactions/'
COMMENT 'Customer transaction history';

-- Create view with row-level security
CREATE VIEW production.customer_data.customers_masked AS
SELECT
  customer_id,
  name,
  CASE
    WHEN IS_MEMBER('pii-access') THEN email
    ELSE 'REDACTED'
  END AS email,
  created_at
FROM production.customer_data.customers;
```

Step 5: Grant Permissions

```
-- Grant catalog access to data analysts
GRANT USAGE ON CATALOG production TO `data-analysts`;
GRANT USAGE ON SCHEMA production.customer_data TO `data-analysts`;
GRANT SELECT ON TABLE production.customer_data.customers_masked TO `data-analysts`;

-- Grant full access to data engineers
GRANT ALL PRIVILEGES ON CATALOG production TO `data-engineers`;

-- Grant model serving access to applications
GRANT EXECUTE ON FUNCTION production.ml_models.predict_churn TO `app-service-principal`;
```

Advanced Features

Row-Level and Column-Level Security

```
-- Row-level security using IS_MEMBER function
CREATE VIEW sales.regional_data AS
SELECT *
FROM sales.all_sales
WHERE
    region = CURRENT_USER()
    OR IS_MEMBER('sales-managers');

-- Column-level security with dynamic masking
CREATE VIEW customers.protected_view AS
SELECT
    customer_id,
    name,
    CASE
        WHEN IS_MEMBER('pii-viewers') THEN ssn
        ELSE 'XXX-XX-XXXX'
    END AS ssn,
    CASE
        WHEN IS_MEMBER('pii-viewers') THEN email
        ELSE REGEXP_REPLACE(email, '^(.{2}).*(@.*)$', '$1***$2')
    END AS email
FROM customers.raw_data;
```

Data Lineage

Unity Catalog automatically tracks lineage for: - Table-to-table transformations - Notebook and job executions - ML model training and serving - Cross-workspace data flows

```
-- Query lineage information
SELECT * FROM system.access.table_lineage
WHERE target_table_full_name = 'production.gold.customer_360'
ORDER BY event_time DESC;
```


Audit Logging

```
-- Query audit logs for data access
SELECT
  event_time,
  user_identity.email,
  request_params.full_name_arg AS table_accessed,
  request_params.command_text AS query_text
FROM system.access.audit
WHERE action_name = 'getTable'
  AND event_date >= CURRENT_DATE() - INTERVAL 7 DAYS
ORDER BY event_time DESC;

-- Monitor privilege grants
SELECT
  event_time,
  user_identity.email AS granted_by,
  request_params.securable_full_name,
  request_params.principal,
  request_params.privileges
FROM system.access.audit
WHERE action_name = 'grant'
ORDER BY event_time DESC;
```

Delta Sharing

```
-- Create a share for external data sharing
CREATE SHARE customer_analytics_share
COMMENT 'Shared customer analytics for partners';

-- Add tables to the share
ALTER SHARE customer_analytics_share
ADD TABLE production.analytics.customer_metrics;

-- Create recipient
CREATE RECIPIENT partner_company
USING ID 'partner-databricks-account-id'
COMMENT 'Partner organization recipient';

-- Grant access to share
GRANT SELECT ON SHARE customer_analytics_share TO RECIPIENT partner_company;
```

Integration with Databricks Services

Unity Catalog + Databricks Apps

```
# Access Unity Catalog tables from Databricks App
import streamlit as st
from databricks import sql
import os

# Connect using app identity
connection = sql.connect(
    server_hostname=os.getenv("DATABRICKS_HOST"),
    http_path=f"/sql/1.0/warehouses/{os.getenv('WAREHOUSE_ID')}"
)

@st.cache_data
def load_data(catalog, schema, table):
    cursor = connection.cursor()
    cursor.execute(f"SELECT * FROM {catalog}.{schema}.{table}")
    return cursor.fetchall_arrow().to_pandas()

# Load data with Unity Catalog governance
df = load_data("production", "customer_data", "customers_masked")
st.dataframe(df)
```

Unity Catalog + Asset Bundles

```
# databricks.yml - Define Unity Catalog resources in bundles
resources:
  schemas:
    customer_schema:
      catalog_name: ${var.catalog}
      name: customer_data
      comment: "Customer data schema"

  grants:
    analyst access:
      securable_type: "SCHEMA"
      securable_name: "${var.catalog}.customer_data"
      principal: "data-analysts"
      privileges: ["USAGE", "SELECT"]
```

Unity Catalog + ML Models

```
import mlflow
from mlflow import MlflowClient

# Set Unity Catalog as registry
mlflow.set_registry_uri("databricks-uc")

# Register model in Unity Catalog
model_name = "production.ml_models.churn_prediction"
mlflow.register_model(
    model_uri=f"runs:{run_id}/model",
    name=model_name
)

# Set model alias
client = MlflowClient()
client.set_registered_model_alias(
    name=model_name,
    alias="champion",
    version=3
)

# Grant model access
spark.sql(f"""
    GRANT EXECUTE ON FUNCTION {model_name}
    TO `ml-serving-principal`
""")
```

Best Practices

1. Catalog Organization

Environment-Based:

```
dev catalog
├─ schema_a
└─ schema_b

staging_catalog
├─ schema_a
└─ schema_b

production catalog
├─ schema_a
└─ schema_b
```

Domain-Based:

```
marketing_catalog
├── campaigns
└── analytics

finance_catalog
├── transactions
└── reporting

engineering_catalog
├── telemetry
└── logs
```

2. Naming Conventions

- Use lowercase with underscores: `customer_data` , `ml_models`
- Include environment prefix when needed: `prod_customer_data`
- Use descriptive schema names: `customer_analytics` not `schema1`
- Document all objects with meaningful comments

3. Security Principles

- **Principle of Least Privilege:** Grant minimum required permissions
- **Use Groups:** Assign permissions to groups, not individual users
- **Separate Environments:** Maintain strict isolation between dev/staging/prod
- **Regular Audits:** Review permissions and access logs regularly
- **Document Policies:** Maintain clear documentation of access policies

4. Performance Optimization

- Use external locations for large datasets
- Implement table partitioning for query performance
- Create materialized views for frequently accessed aggregations
- Monitor query performance through system tables

Migration Strategies

Upgrading from Hive Metastore

```
-- Sync Hive metastore table to Unity Catalog
CREATE TABLE production.migrated_data.customers
DEEP CLONE hive_metastore.default.customers;

-- Create external table pointing to existing data
CREATE EXTERNAL TABLE production.legacy_data.orders
LOCATION 's3://legacy-bucket/orders/';

-- Gradually migrate by creating views
CREATE VIEW production.transition.customers AS
SELECT * FROM hive_metastore.default.customers;
```

Using UCX (Unity Catalog Migration Tool)

```
# Install UCX
databricks labs install ucx

# Assess current workspace
databricks labs ucx assessment

# Create migration plan
databricks labs ucx create-table-mapping

# Execute migration
databricks labs ucx migrate-tables
```

Monitoring and Governance

Key Metrics to Track

```
-- Table access frequency
SELECT
    request_params.full_name_arg AS table_name,
    COUNT(*) AS access_count,
    COUNT(DISTINCT user_identity.email) AS unique_users
FROM system.access.audit
WHERE action_name = 'getTable'
    AND event_date >= CURRENT_DATE() - INTERVAL 30 DAYS
GROUP BY table_name
ORDER BY access_count DESC;

-- Storage usage by catalog
SELECT
    catalog_name,
    SUM(size_in_bytes) / 1024 / 1024 / 1024 AS size_gb,
    COUNT(*) AS table_count
FROM system.information_schema.tables
GROUP BY catalog_name;

-- Permission grants over time
SELECT
    DATE(event_time) AS grant_date,
    COUNT(*) AS grants_count
FROM system.access.audit
WHERE action_name IN ('grant', 'revoke')
GROUP BY grant_date
ORDER BY grant_date DESC;
```

Troubleshooting Common Issues

Issue: "PERMISSION_DENIED" Errors

```
-- Check current user privileges
SHOW GRANTS ON CATALOG production;
SHOW GRANTS ON SCHEMA production.customer data;
SHOW GRANTS ON TABLE production.customer_data.customers;

-- Verify group membership
SELECT CURRENT_USER();
SELECT * FROM system.access.group_membership
WHERE user_name = CURRENT_USER();
```

Issue: Cannot Access External Tables

```
-- Verify storage credential
DESCRIBE STORAGE CREDENTIAL aws_s3_credential;

-- Check external location
DESCRIBE EXTERNAL LOCATION production_data;

-- Test access
SELECT * FROM production.external_data.test_table LIMIT 10;
```

Issue: Lineage Not Appearing

- Ensure compute has Unity Catalog enabled
- Verify table is registered in Unity Catalog (not Hive metastore)
- Check that lineage capture is enabled in workspace settings
- Allow 24-48 hours for lineage to populate

References

- [Official Unity Catalog Documentation](#)
 - [Unity Catalog Best Practices](#)
 - [Unity Catalog GitHub Repository](#)
 - [Delta Sharing Protocol](#)
 - [UCX Migration Tool](#)
-